Prathamesh More

Maharashtra, Pune | pratmore@deloitte.com | +91 8482847824

Summary

Data Scientist with 4+ years of experience in AI/ML and Generative AI, specializing in Agentic AI systems and Retrieval Augmented Generation (RAG). Hands-on experience building multi-agent architectures with autonomous orchestration, hybrid retrieval pipelines, and schema-driven answer generation grounded strictly in retrieved context. Proficient in Python, LangChain, LangGraph, and vector databases.

Education

Bachelors of Engineering (Mechanical) - JSPM's Rajarshi Shahu College of Engineering — CGPA: 8.51

Skills

- **Programming Languages:** Python, Java, C
- **Machine Learning Algorithms:** Logistic Regression, Linear Regression, SVMs, Decision Trees, KNN, Ensembles
- **Deep Learning:** ANN, RNN, LSTM, CNN
- **Gen AI Frameworks:** Langchain, Langgraph, Deep Agents
- **Cloud:** Azure
- **SQL:** MySQL, Oracle, Postgres
- **NoSQL:** MongoDB
- **Frameworks:** Pandas, Numpy, Matplotlib, Scikit-learn, Keras, Tensorflow
- **NLP:** Tokenization, Chunking, Embedding
- **LLMs:** GPT-4o, gpt-oss-120b, Llama3.3-70b

Experience

Consultant, **Deloitte USI** — May 2025 - Present

- Building a scalable **Multi-Agent** platform using the **DeepAgents** framework (LangGraph-based), where LLM agents autonomously discover skills via SKILL.md instructions, execute **Text2SQL** and **RAG** workflows through a middleware-driven tool-calling architecture, and dynamically classify and route queries — replacing hard-coded conditional edges with LLM-driven orchestration.
- Designed a **two-stage RAG pipeline** with hybrid retrieval (BM25 + dense vectors via **Milvus**), **Reciprocal Rank Fusion**, **Cohere reranking**, and source-balanced chunk selection — paired with a document ingestion pipeline using **Docling** and **PyMuPDF** parsers to handle heterogeneous PDFs with contextual chunking and domain-specific metadata enrichment.
- Built a production-resilient **middleware stack** — circuit breakers, rate limiters, protocol parsers for reasoning-model quirks, and agent flow loggers — integrated with **OpenAI** embeddings, **gpt-oss-120b** for complex reasoning, and **PostgreSQL** as the transactional SQL backend.
- Developed schema-driven **answer generation** with YAML-based formatting rules, answer-type routing (comparison, causal, procedural, factual), and inline citation enforcement — grounding every response strictly in retrieved context to eliminate hallucination.

Data Scientist, **Tech Mahindra** — Nov 2024 – May 2025

- Developed an enterprise-grade **RAG**-based chatbot capable of retrieving and synthesizing information from 4,000+ heterogeneous PDFs containing structured and unstructured content.
- Built secure, modular pipelines using **LangChain** and **Azure** services, incorporating authentication, access control, and environment-specific configurations.
- Utilized **OpenAI** embedding models with **Chroma** Vector DB and **GPT-4o**; evaluated system quality using relevancy scoring between ground truth answers and generated responses.
- Implemented Agentic RAG with query decomposition, context validation, and iterative retrieval, improving answer accuracy by 4%.
- Added multilingual query and response support, increasing user adoption and traffic by 11
- Optimized chunking strategies and retrieval parameters to balance latency, context quality, and token efficiency.

AI/ML Developer, **Tata Consultancy Services** — Oct 2021 – Oct 2024

- Developed and deployed predictive ML models including a **Policy Compliance Tracker** and **Card Limit Optimizer** for enterprise financial workflows.
- Performed end-to-end data preprocessing, feature engineering, and exploratory analysis to improve data quality and model robustness.
- Automated compliance monitoring workflows, improving detection accuracy by 16% and reducing operational monitoring costs by 18%.
- Collaborated closely with business stakeholders to translate requirements into ML solutions, iteratively refining models based on real-world feedback.
- Enhanced end-user experience and system usability, contributing to an 8% increase in positive feedback and up to 12% cost savings through process optimization.

Certifications

- Generative AI with Langchain and Huggingface
- Microsoft Azure AI-900
- Deep Learning (NIELIT)